

Ch6 – Final Assignment

Tesla stock price prediction using Times Series Prediction, Deep learning



Wang Xu finish at 2022-4-28

Friendly remind: This report only shows a part of code.

Please view whole code in my github link : [ipynb link](#)



Contents ⚙

- 1 Assignment - Tesla stack price predict - ARIMA, SARIMA, LSTM
- 2 Main objective
- 3 Data description
- ▼ 4 Exploration Data Analysis
 - 4.1 Import libraries
 - 4.2 Data cleaning
 - 4.3 Visualization
 - ▼ 4.4 Check its stationarity
 - 4.4.1 split chunks to verify mean, variance
 - 4.4.2 Whether is normal distribution
 - 4.4.3 Statistics - AdFuller test
- ▼ 5 Model preparing
 - 5.1 Split train(2010:2015) and test dataset
 - 5.2 Differencing - transform to stationarity
 - 5.3 ACF, PACF for q,p
- ▼ 6 ARIMA model
 - 6.1 Parameters choosing from acf, pacf, order=(1,1,1)
 - 6.2 Prediction using model_1
 - 6.3 Validation : Mse
- ▼ 7 Pmdarima Model
 - 7.1 Parameters from pmd, order=(1,0,1), seasonal=False
 - 7.2 Prediction using model_2_best
 - 7.3 Validation: mse
 - 7.4 Short Summary
- ▼ 8 SARIMA model
 - 8.1 Parameters, best order, seasonal=True
 - 8.2 Prediction using model_3_best
 - 8.3 Validation: mse
- 9 Summary - key findings - Models comparing
- 10 Insights and Suggestions
- ▼ 11 Further revisiting (experiment)
 - ▼ 11.1 Change train data (2013:2016)
 - 11.1.1 transfer to stationarity by differencing
 - 11.1.2 Choosing parameters
 - 11.1.3 SARIMA model try
 - 11.1.4 Prediction with new train, new parameters
 - ▼ 11.2 Deep learning model trying- LSTM
 - 11.2.1 Model preparing
 - 11.2.2 split train-test set
 - 11.2.3 LSTM Model
- 12 Final conclusion:



Objectives

- The main objectives of this assignment is focusing on **time series prediction** of Tesla stock price and choose a preferred model.
- Experiments specifically on parameter variations of models:
 - **ARIMA** (Autoregression Integrated Moving Average);
 - **SARIMA** (Seasonal Autoregression Integrated Moving Average) ;
 - and **deep learning LSTM** (Long-Short Term Memory Network).
- It may give benefits as advices or price trend predictions for individual investors, investment funds, any other stock research organizations, etc.

Brief Description of dataset

- Tesla's stock price data from its initial public offering (IPO) date 2010-6-29 to 2017-3-17. (1692 rows)
- Data download from: <https://www.kaggle.com/datasets/rpaguirre/tesla-stock-price>

7 Attributes:

1. "Date" - The date
 2. "Open" - The opening price of the stock
 3. "High" - The high price of that day
 4. "Low" - The low price of that day
 5. "Close" -The closed price of that day
 6. "Volume" - The amount of stocks traded during that day
 7. "Adjusted Close" - The stock's closing price that has been amended to include any distributions/corporate actions that occurs before next days open
- **In this case, "Close" price will be considered as main prediction.**

4.2 Data cleaning

```
: 1 df['Date'] = pd.to_datetime(df['Date'])
: 2 df.set_index('Date', inplace=True)
```

```
: 1 df.head()
```

| | | Open | High | Low | Close | Volume | Adj Close |
|---|------------|------|------|------|-------|----------|-----------|
| | Date | | | | | | |
| 1 | 2010-06-29 | 19.0 | 25.0 | 17.5 | 23.9 | 18766300 | 23.9 |
| 2 | 2010-06-30 | 25.8 | 30.4 | 23.3 | 23.8 | 17187100 | 23.8 |
| 3 | 2010-07-01 | 25.0 | 25.9 | 20.3 | 22.0 | 8218800 | 22.0 |
| 4 | 2010-07-02 | 23.0 | 23.1 | 18.7 | 19.2 | 5139800 | 19.2 |
| 5 | 2010-07-06 | 20.0 | 20.0 | 15.8 | 16.1 | 6866900 | 16.1 |

```
: 1 df.describe()
```

| | Open | High | Low | Close | Volume | Adj Close |
|-------|---------|---------|---------|---------|--------------|-----------|
| count | 1,692.0 | 1,692.0 | 1,692.0 | 1,692.0 | 1,692.0 | 1,692.0 |
| mean | 132.4 | 134.8 | 130.0 | 132.4 | 4,270,740.9 | 132.4 |
| std | 94.3 | 95.7 | 92.9 | 94.3 | 4,295,971.3 | 94.3 |
| min | 16.1 | 16.6 | 15.0 | 15.8 | 118,500.0 | 15.8 |
| 25% | 30.0 | 30.6 | 29.2 | 29.9 | 1,194,350.0 | 29.9 |
| 50% | 156.3 | 162.4 | 153.2 | 158.2 | 3,180,700.0 | 158.2 |
| 75% | 220.6 | 224.1 | 217.1 | 220.0 | 5,662,100.0 | 220.0 |
| max | 287.7 | 291.4 | 280.4 | 286.0 | 37,163,900.0 | 286.0 |

4.1 Import libraries

```
: 1 # imports
: 2 import pandas as pd
: 3 import numpy as np
: 4 import matplotlib.pyplot as plt
: 5 import seaborn as sns
: 6 import warnings
: 7 from datetime import datetime
: 8 import os
: 9 os.chdir('data')
:10 from colorsetup import colors, palette
:11 sns.set_palette(palette)
:12 # ignore warnings
:13 warnings.filterwarnings('ignore')
:14 pd.options.display.float_format = '{:.1f}'.format
:15 %matplotlib inline
:16 plotsize = (12, 8)
:17 from statsmodels.tsa.stattools import adfuller
```

```
: 1 df = pd.read_csv("Tesla_stock.csv")
```

```
: 1 df.head()
```

| | Date | Open | High | Low | Close | Volume | Adj Close |
|---|-----------|------|------|------|-------|----------|-----------|
| 0 | 6/29/2010 | 19.0 | 25.0 | 17.5 | 23.9 | 18766300 | 23.9 |
| 1 | 6/30/2010 | 25.8 | 30.4 | 23.3 | 23.8 | 17187100 | 23.8 |
| 2 | 7/1/2010 | 25.0 | 25.9 | 20.3 | 22.0 | 8218800 | 22.0 |
| 3 | 7/2/2010 | 23.0 | 23.1 | 18.7 | 19.2 | 5139800 | 19.2 |
| 4 | 7/6/2010 | 20.0 | 20.0 | 15.8 | 16.1 | 6866900 | 16.1 |



4.3 Visualization

```
1 plt.figure(figsize=(18,8))
2 plt.plot(df['Close'], 'b-')
3 plt.title("Tesla - Close price")
4 plt.ylabel("Stock price")
5 plt.grid()
```



Check its stationarity

- The prediction should be more reasonable on the close price.
- from close price plot, it seems that is non-stationarity.

4.4.1 split chunks to verify mean, variance

```
1 close_array = pd.Series(df['Close'].values)
2 chunks = np.split(close_array, indices_or_sections=9)
3 print("{}|{:7}|{}".format("chunks","mean","variance"))
4 print("-*25)
5 for i, j in enumerate(chunks,1):
6     print("{:6}|{:.6}|{:.6}".format(i, np.mean(j), np.var(j)))
```

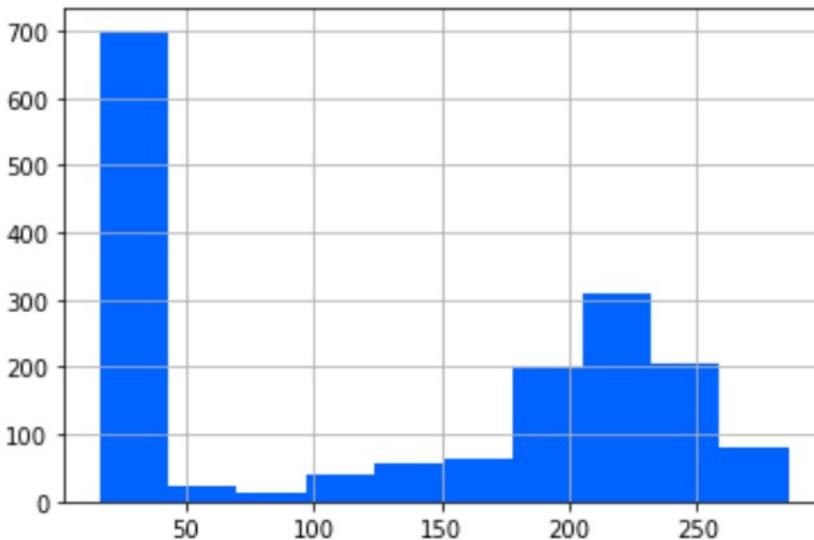
| chunks | mean | variance |
|--------|---------|----------|
| 1 | 23.6027 | 18.6438 |
| 2 | 27.5632 | 7.33095 |
| 3 | 31.0821 | 7.76437 |
| 4 | 46.2839 | 539.716 |
| 5 | 164.471 | 1263.25 |
| 6 | 230.328 | 557.292 |
| 7 | 230.528 | 680.763 |
| 8 | 218.87 | 580.545 |
| 9 | 219.128 | 596.781 |

It has changing means and changing variances.

4.4.2 Whether is normal distribution

```
1 close_array.hist()
```

<AxesSubplot:>



It is not normal distribution.

4.4.3 Statistics - AdFuller test

Addiction to visualization, statistics proving whether it is stationarity helps as well.

```
: 1 adf, pvalue, usedlag, nobs, critical_values, icbest = adfuller(df['Close'])
 2 print("Adf:",adf)
 3 print("p-value:",pvalue)
```

```
Adf: -0.8137329693176628
p-value: 0.8150547908640395
```

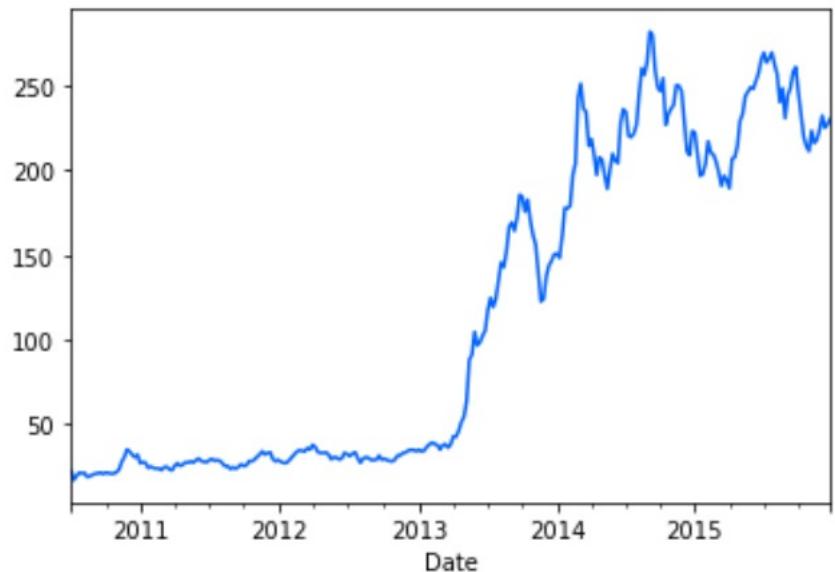
p-value > 0.05 clearly, it is non-stationarity.

Non-stationarity!

Model preparing

5.1 Split train(2010:2015) and test dataset

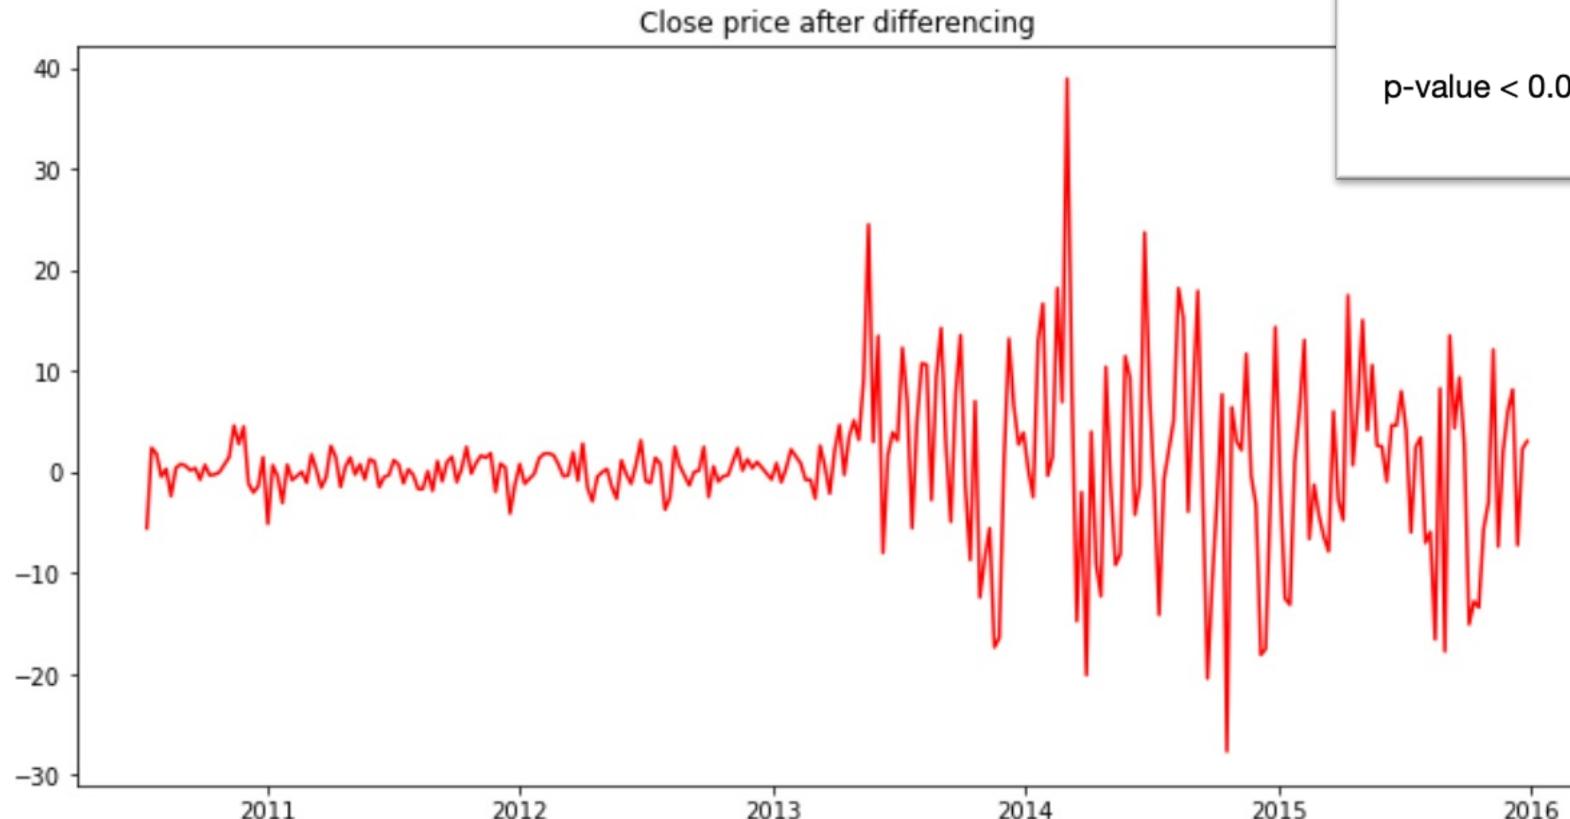
```
1 close_week = df['Close'].resample('W').mean()  
1 stock_train = close_week['2010':'2015']  
1 stock_train.plot()  
<AxesSubplot:xlabel='Date'>
```



Choose train-set :
2010 – 2015 close
price in weeks
&
Transfer to stationarity

5.2 Differencing - transform to stationarity

```
: 1 stock_diff = stock_train.diff()  
2 stock_diff = stock_diff.dropna()  
3 plt.figure(figsize=(12,6))  
4 plt.plot(stock_diff,'r-')  
5 plt.title("Close price after differencing")  
:  
Text(0.5, 1.0, 'Close price after differencing')
```



```
1 adfuller(stock_diff)
```

```
(-13.684077570268537,  
 1.3829853031786577e-25,  
 0,  
 285,  
 {'1%': -3.4535050041524245,  
 '5%': -2.8717352599720294,  
 '10%': -2.5722024776854417},  
 1841.0821374221925)
```

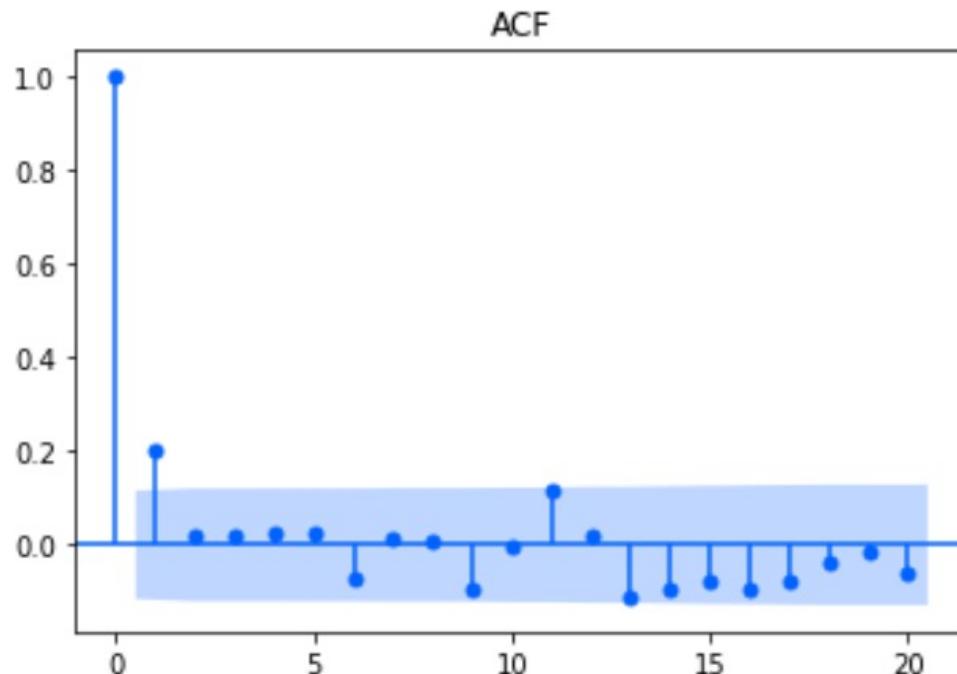
p-value < 0.05, now after differencing, the stock_diff is stationary

Become more
stationarity!

Parameters choosing from ACF (autocorrelation plot), PACF (partial autocorrelation plot) in (p,q)

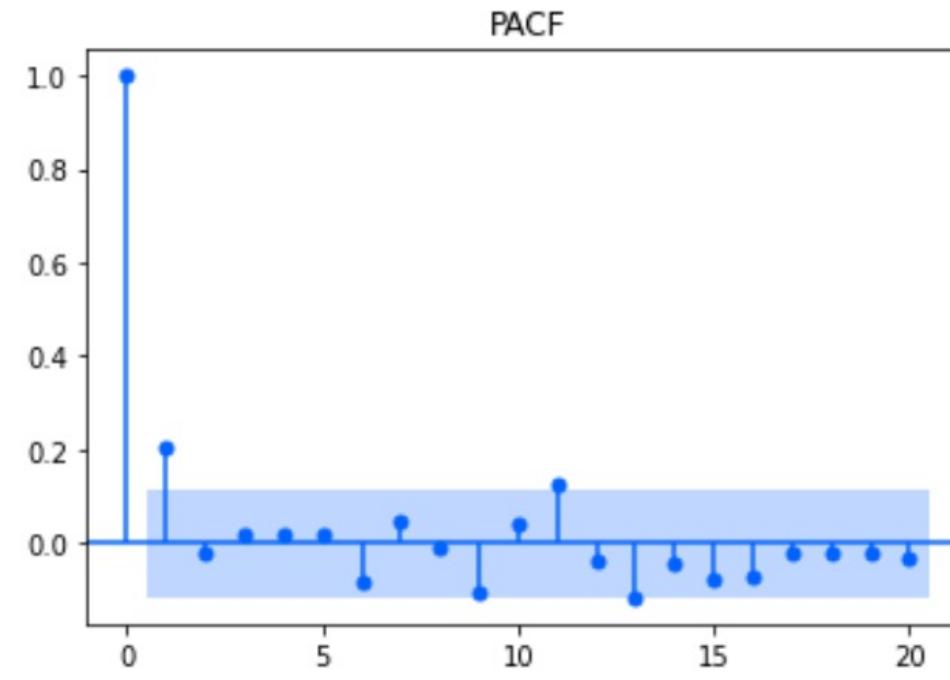
```
1 plot_acf(stock_diff, lags=20)
2 plt.title("ACF")
```

Text(0.5, 1.0, 'ACF')



```
1 plot_pacf(stock_diff, lags=20)
2 plt.title("PACF")
```

Text(0.5, 1.0, 'PACF')



From above plots, it seems that q=1, p=1 is preferred.

Baseline: Model_1, ARIMA, order=(1,1,1)

Information criteria (AIC, BIC)
which penalize the number of
parameters the model uses

6.1 Parameters choosing from acf, pacf, order=(1,1,1)

```
1 from statsmodels.tsa.arima.model import ARIMA
1 model_1 = ARIMA(stock_train,order=(1,1,1),freq='W').fit()
1 print(model_1.summary())
```

SARIMAX Results

| Dep. Variable: | Close | No. Observations: | 287 |
|------------------|-------------------------|-------------------|----------|
| Model: | ARIMA(1, 1, 1) | Log Likelihood | -969.345 |
| Date: | Thu, 28 Apr 2022 | AIC | 1944.690 |
| Time: | 10:37:50 | BIC | 1955.658 |
| Sample: | 07-04-2010 - 12-27-2015 | HQIC | 1949.086 |
| Covariance Type: | opg | | |

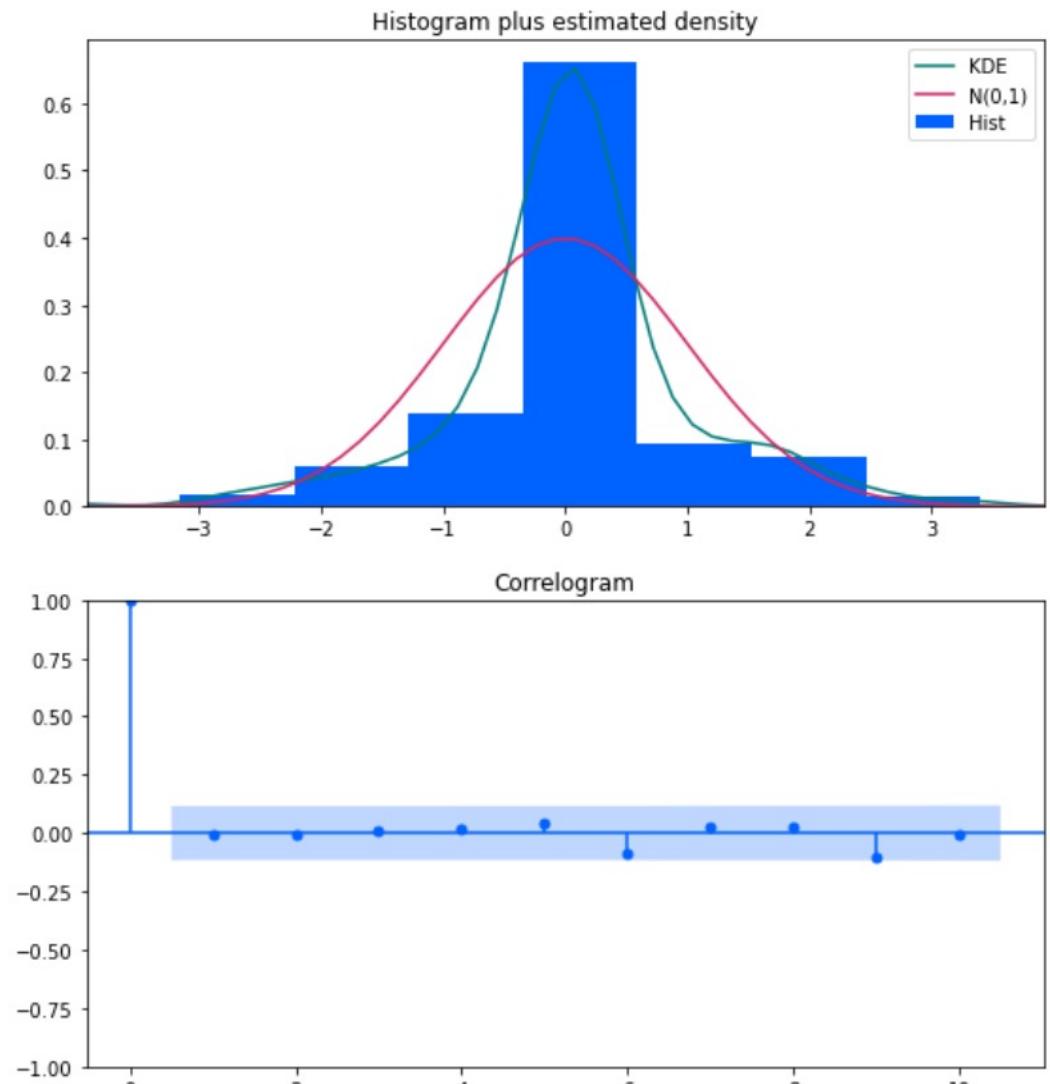
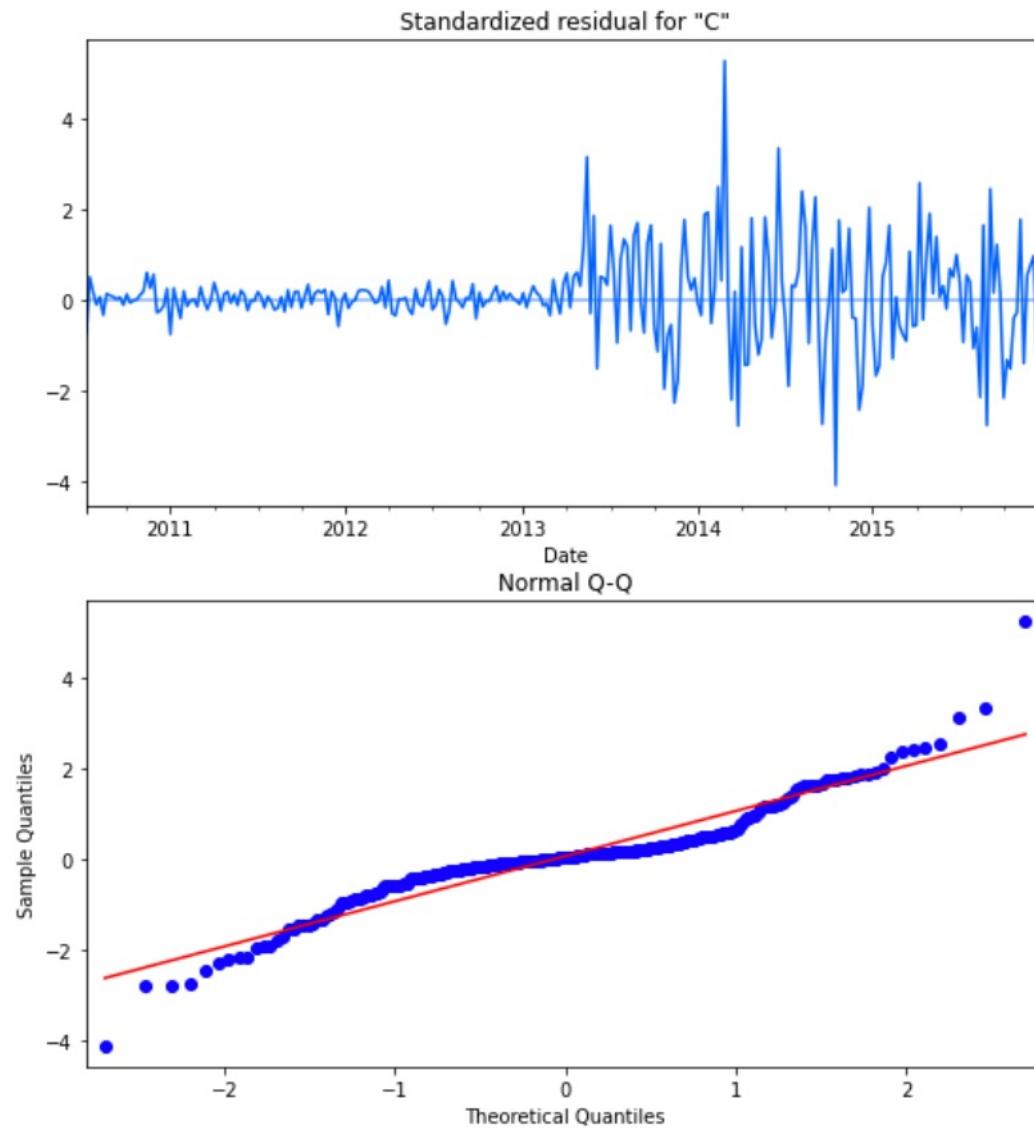
| | coef | std err | z | P> z | [0.025 | 0.975] |
|--------|---------|---------|--------|-------|--------|--------|
| ar.L1 | 0.1321 | 0.190 | 0.697 | 0.486 | -0.240 | 0.504 |
| ma.L1 | 0.0840 | 0.187 | 0.451 | 0.652 | -0.282 | 0.450 |
| sigma2 | 51.4491 | 2.571 | 20.014 | 0.000 | 46.411 | 56.487 |

| Ljung-Box (L1) (Q): | 0.01 | Jarque-Bera (JB): | 213.03 |
|-------------------------|-------|-------------------|--------|
| Prob(Q): | 0.91 | Prob(JB): | 0.00 |
| Heteroskedasticity (H): | 33.60 | Skew: | 0.30 |
| Prob(H) (two-sided): | 0.00 | Kurtosis: | 7.19 |

Warnings:

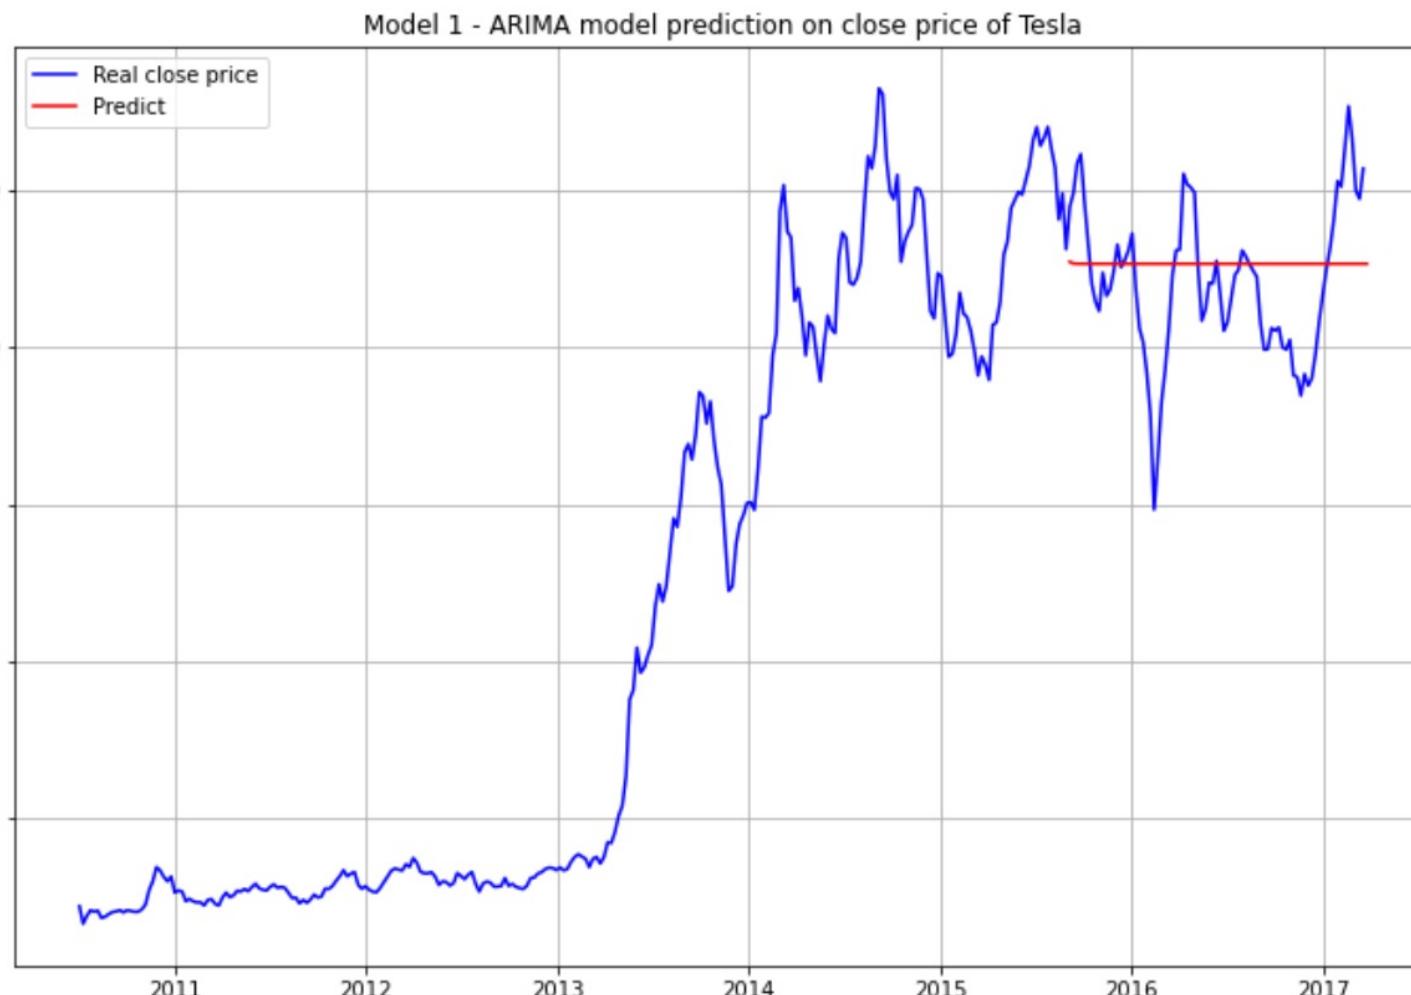
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
1 model_1.plot_diagnostics(figsize=(20,10));
```



6.2 Prediction using model_1

```
1 predict_1 = model_1.predict(start=270,end=351,dynamic=True)
2 plt.figure(figsize=(12,8))
3 plt.plot(close_week,"b-",label="Real close price")
4 plt.plot(predict_1,'r-',label='Predict')
5 plt.legend(loc='best')
6 plt.title("Model 1 - ARIMA model prediction on close price of Tesla")
7 plt.grid(0.3)
```



6.3 Validation : Mse

```
1 num_weeks_test = len(predict_1)
2 close_test = close_week[-num_weeks_test:] #find the last n weeks of data
```



```
1 def mse(observation, estimates):
2     assert type(observation)==type(np.array([]))
3     assert type(estimates)==type(np.array([]))
4     assert len(observation)==len(estimates)
5     diff = abs(observation) - abs(estimates)
6     square = diff **2
7     mse = sum(square)
8     return mse
```

```
1 mse(close_test.values, predict_1.values)
```

50246.62748995037

Model_2, ARIMA, order=(1,0,1)

```
1 model_2 = pmd.auto_arima(stock_train,start_p=0,
2                             start_q=0,max_p=3,max_q=3,
3                             m=11, d=0, D=1,trace=True,
4                             seasonal=False,
5                             error_action='ignore',
6                             suppress_warnings=True,stepwise=True)
```

```
Performing stepwise search to minimize aic
ARIMA(0,0,0)(0,0,0)[0] : AIC=3680.755, Time=0.01 sec
ARIMA(1,0,0)(0,0,0)[0] : AIC=inf, Time=0.10 sec
ARIMA(0,0,1)(0,0,0)[0] : AIC=inf, Time=0.13 sec
ARIMA(1,0,1)(0,0,0)[0] : AIC=1957.680, Time=0.11 sec
ARIMA(2,0,1)(0,0,0)[0] : AIC=1959.475, Time=0.15 sec
ARIMA(1,0,2)(0,0,0)[0] : AIC=1959.510, Time=0.10 sec
ARIMA(0,0,2)(0,0,0)[0] : AIC=2994.108, Time=0.13 sec
ARIMA(2,0,0)(0,0,0)[0] : AIC=inf, Time=0.11 sec
ARIMA(2,0,2)(0,0,0)[0] : AIC=inf, Time=0.33 sec
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=1958.405, Time=0.18 sec
```

Best model: ARIMA(1,0,1)(0,0,0)[0]
Total fit time: 1.365 seconds

SARIMAX Results

```
=====
Dep. Variable: Close No. Observations: 287
Model: ARIMA(1, 0, 1) Log Likelihood -975.209
Date: Thu, 28 Apr 2022 AIC 1958.417
Time: 14:22:46 BIC 1973.055
Sample: 07-04-2010 HQIC 1964.284
- 12-27-2015
Covariance Type: opg
=====
```

| | coef | std err | z | P> z | [0.025 | 0.975] |
|--------|----------|---------|---------|-------|---------|---------|
| const | 113.4013 | 105.105 | 1.079 | 0.281 | -92.602 | 319.404 |
| ar.L1 | 0.9961 | 0.006 | 161.999 | 0.000 | 0.984 | 1.008 |
| ma.L1 | 0.2132 | 0.047 | 4.553 | 0.000 | 0.121 | 0.305 |
| sigma2 | 51.3949 | 2.440 | 21.065 | 0.000 | 46.613 | 56.177 |

```
=====
```

| | | | |
|-------------------------|-------|-------------------|--------|
| Ljung-Box (L1) (Q): | 0.00 | Jarque-Bera (JB): | 228.76 |
| Prob(Q): | 0.97 | Prob(JB): | 0.00 |
| Heteroskedasticity (H): | 33.44 | Skew: | 0.43 |
| Prob(H) (two-sided): | 0.00 | Kurtosis: | 7.29 |

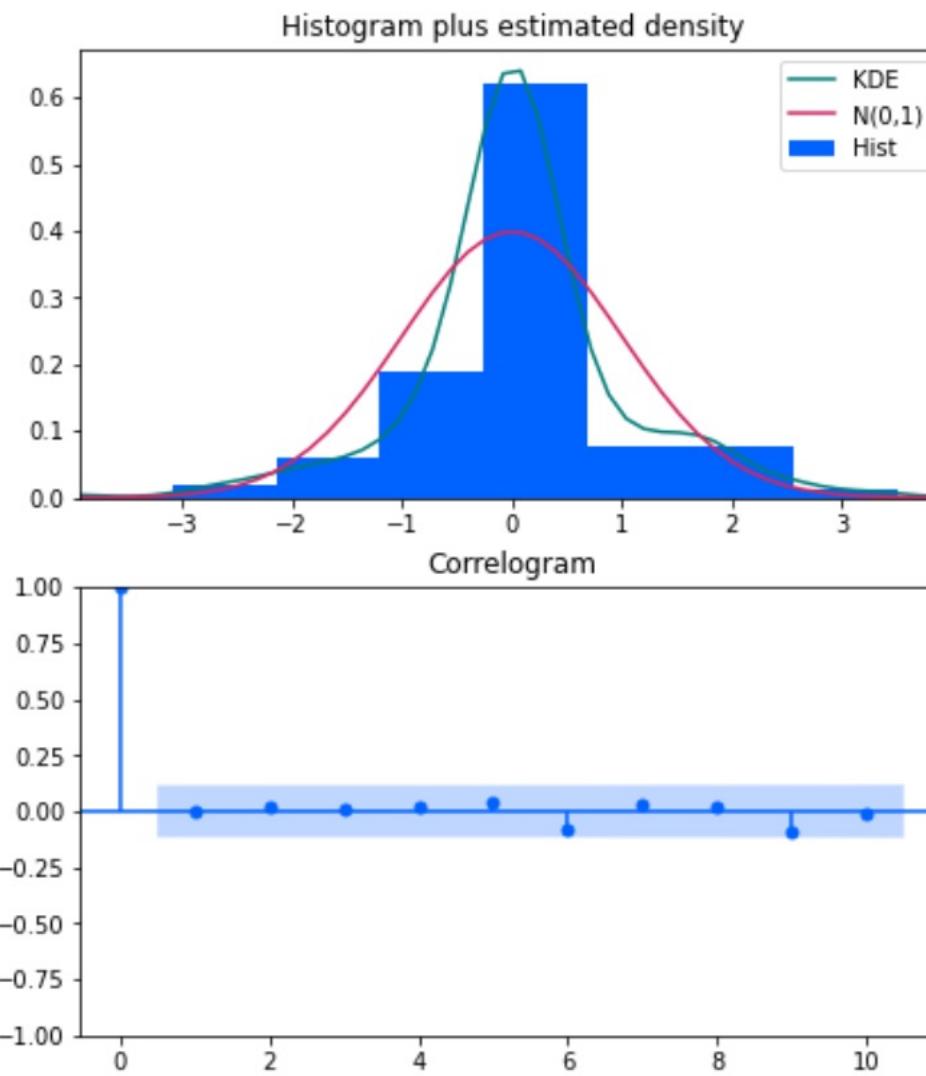
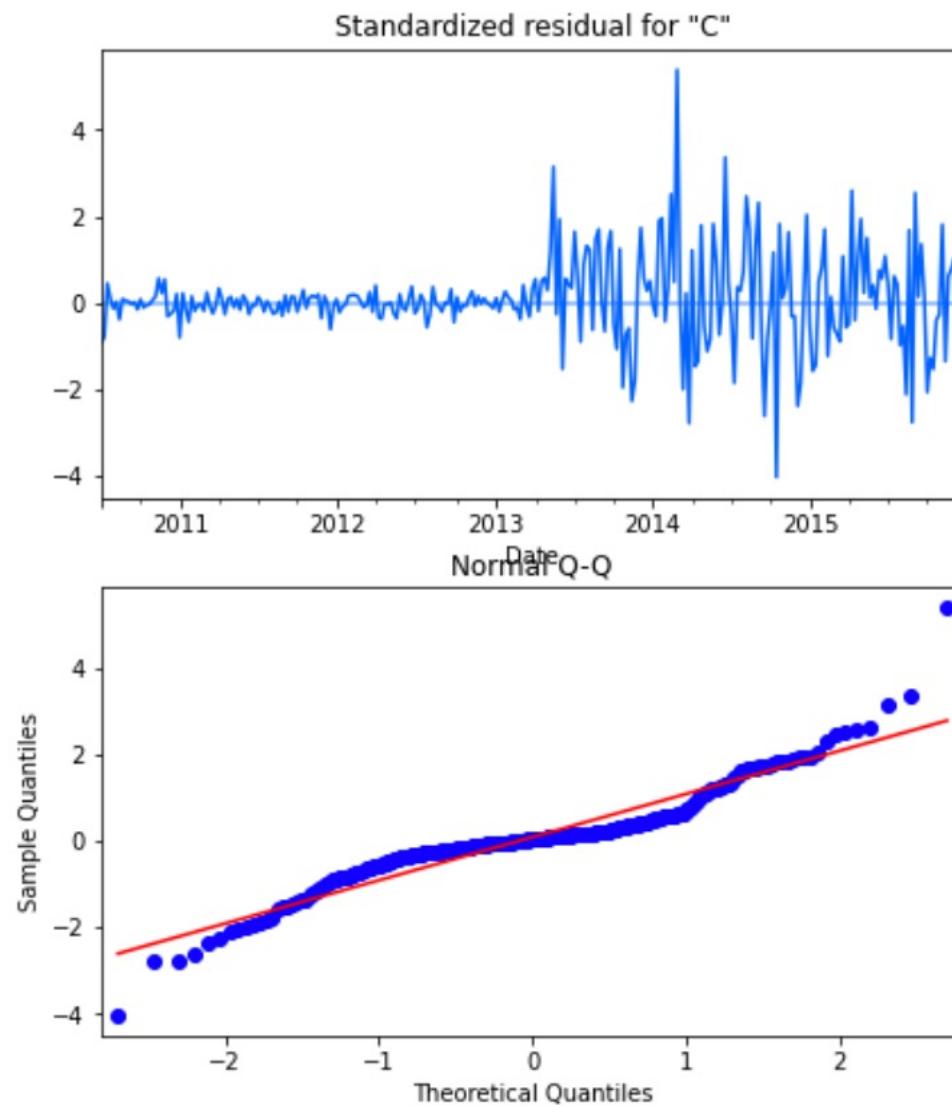
```
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

- Using pmd module to find best parameters
- Information criteria (AIC, BIC) increased.
(penalize the number of parameters the model uses)

```
1 model_2_best.plot_diagnostics(figsize=(15,8));
```



7.2 Prediction using model_2_best

```
1 predict_2 = model_2_best.predict(start=270,end=351,dynamic=True)
2 plt.figure(figsize=(12,8))
3 plt.plot(close_week,"b-",label="Real close price")
4 plt.plot(predict_2,'r-',label='Predict')
5 plt.legend(loc='best')
6 plt.title("Model_2_best - pmd-ARIMA model prediction on close price of Tesla")
7 plt.grid(0.3)
```

Model_2_best - pmd-ARIMA model prediction on close price of Tesla



Mean square error

7.3 Validation: mse

```
1 print("MSE_model_1 (order = (1,1,1)):", mse(close_test.values,pred)
2 print("MSE_model_2_best (order = (1,0,1)):",mse(close_test.values,
```

MSE_model_1 (order = (1,1,1)): 50246.62748995037
MSE_model_2_best (order = (1,0,1)): 61624.472763307844

from the comparison of model_1 and model_2, and their mse:

- Model_1's trend is more accurate than model_2.
- In this case, model_1 is preferred.

Model_3, SARIMA, order=(2,0,1), Seasonal order=(2,1,0,11)

- Using pmd module to find best parameters
- Similar information criteria (AIC, BIC)

```
1 print("best order:", model_3.order)
2 print("best seasonal order:",model_3.seasonal_order)

best order: (2, 0, 1)
best seasonal order: (2, 1, 0, 11)
```

```
1 model_3_best = sm.tsa.statespace.SARIMAX(stock_train,
2                                         order=(2,0,1),
3                                         seasonal_order=(2,1,0,11),
4                                         trend='c').fit()
```

```
1 #comparing model_2 order=(1,0,1) vs. model_3 adding seasonal effect
2 print(model_3_best.summary())
```

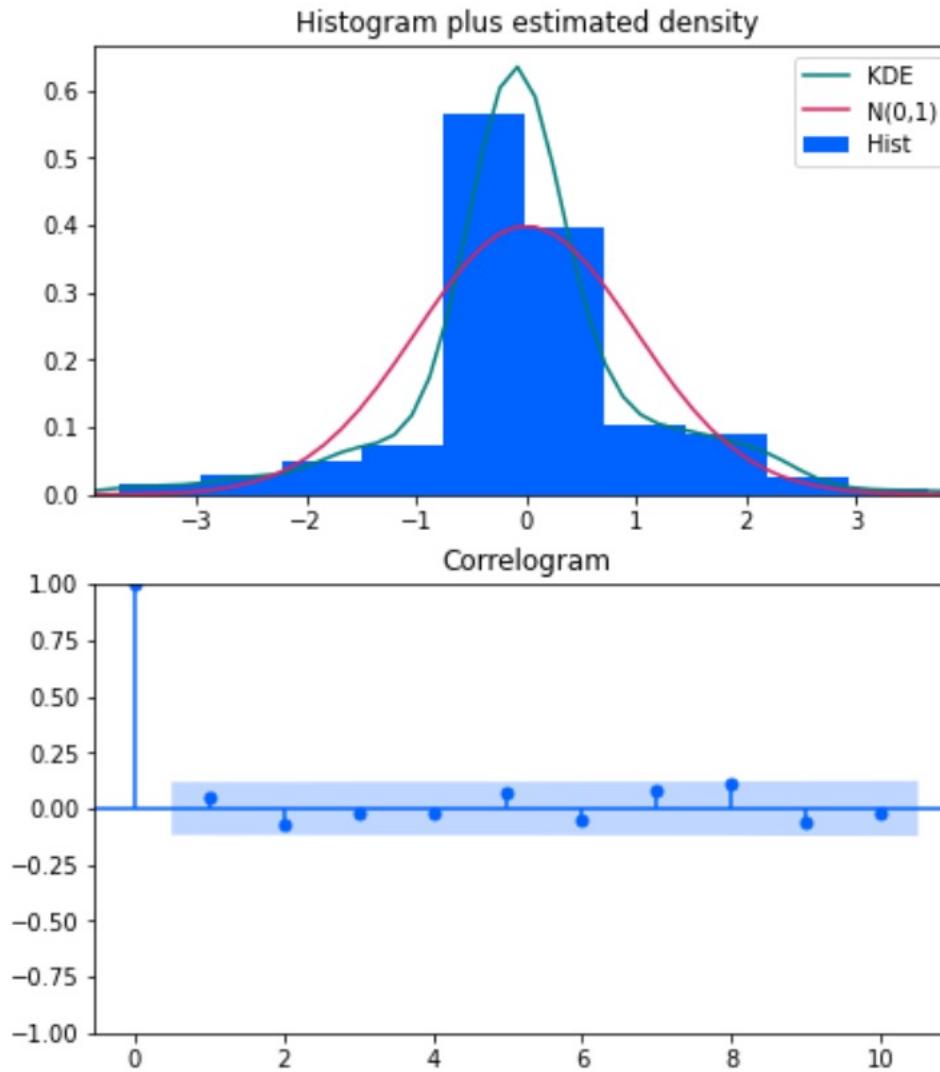
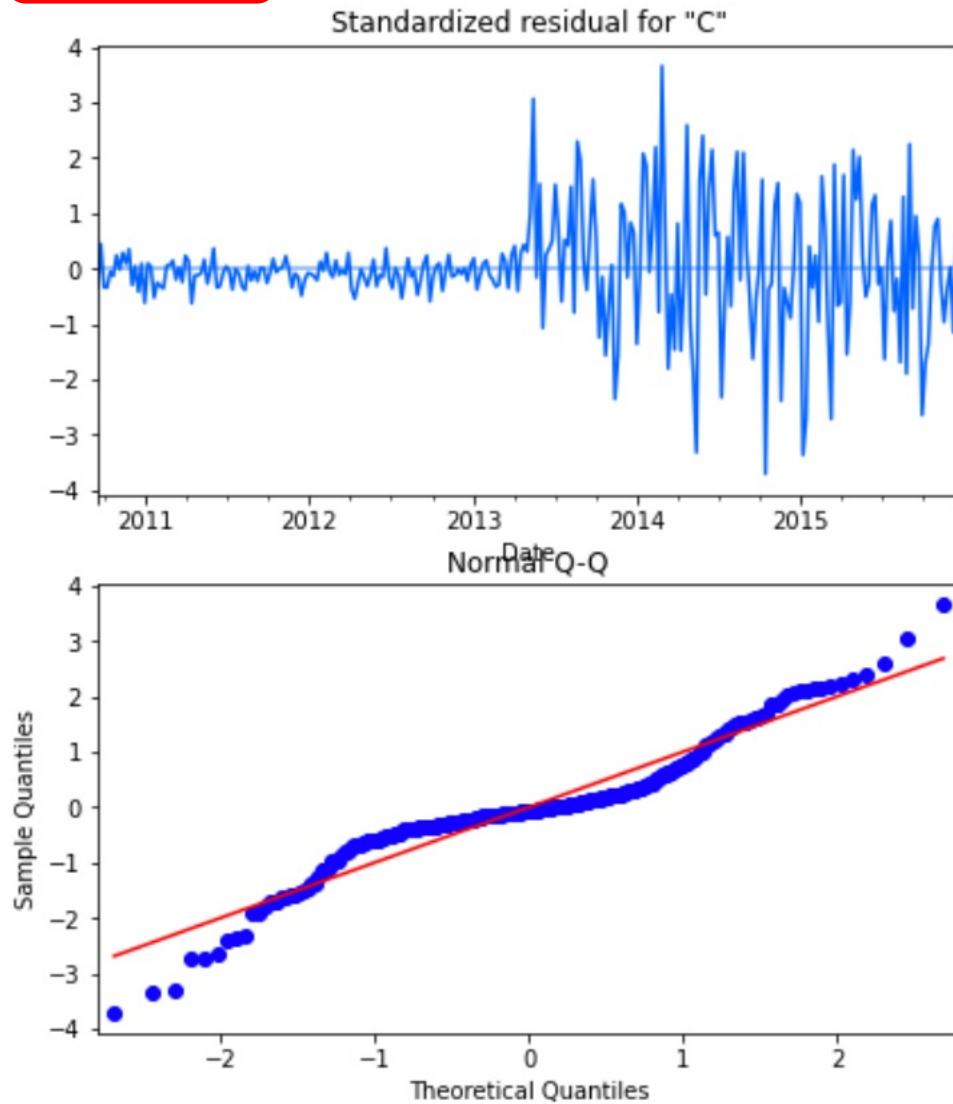
SARIMAX Results

| Dep. Variable: | Close | No. Observations: | 287 | | | |
|-------------------------|---------------------------------|-------------------|----------|-------|--------|--------|
| Model: | SARIMAX(2, 0, 1)x(2, 1, [], 11) | Log Likelihood | -969.047 | | | |
| Date: | Thu, 28 Apr 2022 | AIC | 1952.094 | | | |
| Time: | 23:08:26 | BIC | 1977.437 | | | |
| Sample: | 07-04-2010 - 12-27-2015 | HQIC | 1962.264 | | | |
| Covariance Type: | opg | | | | | |
| | coef | std err | z | P> z | [0.025 | 0.975] |
| intercept | 0.2764 | 0.196 | 1.413 | 0.158 | -0.107 | 0.660 |
| ar.L1 | 1.7794 | 0.097 | 18.384 | 0.000 | 1.590 | 1.969 |
| ar.L2 | -0.7998 | 0.091 | -8.771 | 0.000 | -0.979 | -0.621 |
| ma.L1 | -0.6886 | 0.124 | -5.571 | 0.000 | -0.931 | -0.446 |
| ar.S.L11 | -0.5161 | 0.050 | -10.286 | 0.000 | -0.614 | -0.418 |
| ar.S.L22 | -0.1221 | 0.040 | -3.018 | 0.003 | -0.201 | -0.043 |
| sigma2 | 64.3643 | 4.006 | 16.065 | 0.000 | 56.512 | 72.217 |
| Ljung-Box (L1) (Q): | 0.80 | Jarque-Bera (JB): | 62.66 | | | |
| Prob(Q): | 0.37 | Prob(JB): | 0.00 | | | |
| Heteroskedasticity (H): | 34.49 | Skew: | -0.05 | | | |
| Prob(H) (two-sided): | 0.00 | Kurtosis: | 5.33 | | | |

Warnings:

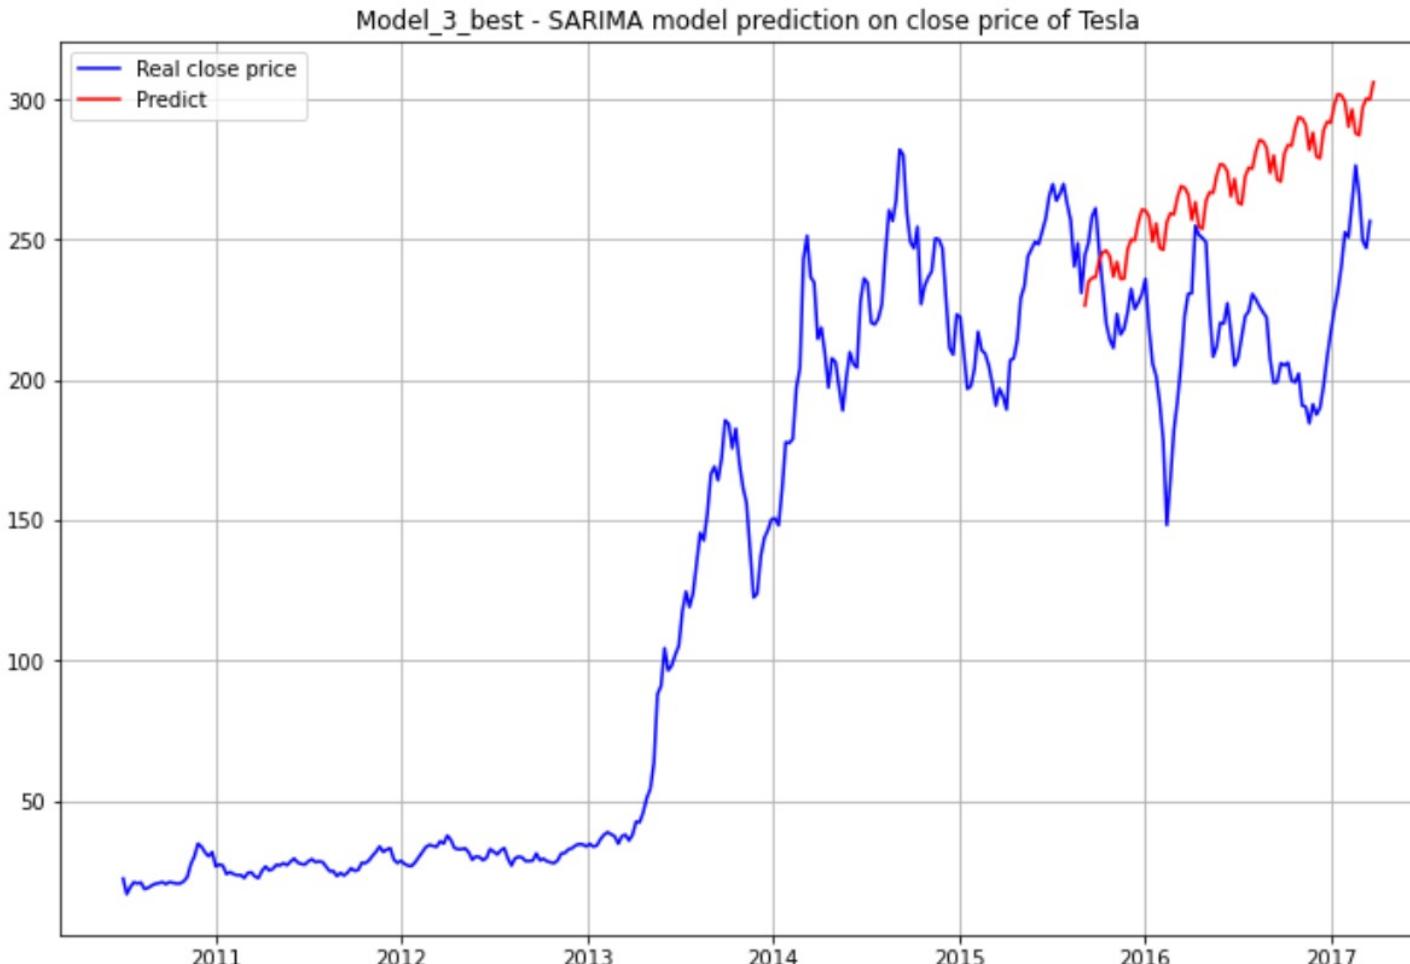
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
1 model_3_best.plot_diagnostics(figsize=(15,8));
```



8.2 Prediction using model_3_best

```
1 predict_3 = model_3_best.predict(start=270,end=351,dynamic=True)
2 plt.figure(figsize=(12,8))
3 plt.plot(close_week,'b-',label="Real close price")
4 plt.plot(predict_3,'r-',label='Predict')
5 plt.legend(loc='best')
6 plt.title("Model_3_best - SARIMA model prediction on close price of Tesla")
7 plt.grid(0.3)
```



8.3 Validation: mse

```
1 print("MSE of model_1:", mse(close_test))
2 print("MSE of model_2_best:", mse(close))
3 print("MSE of model_3_best:", mse(close))
```

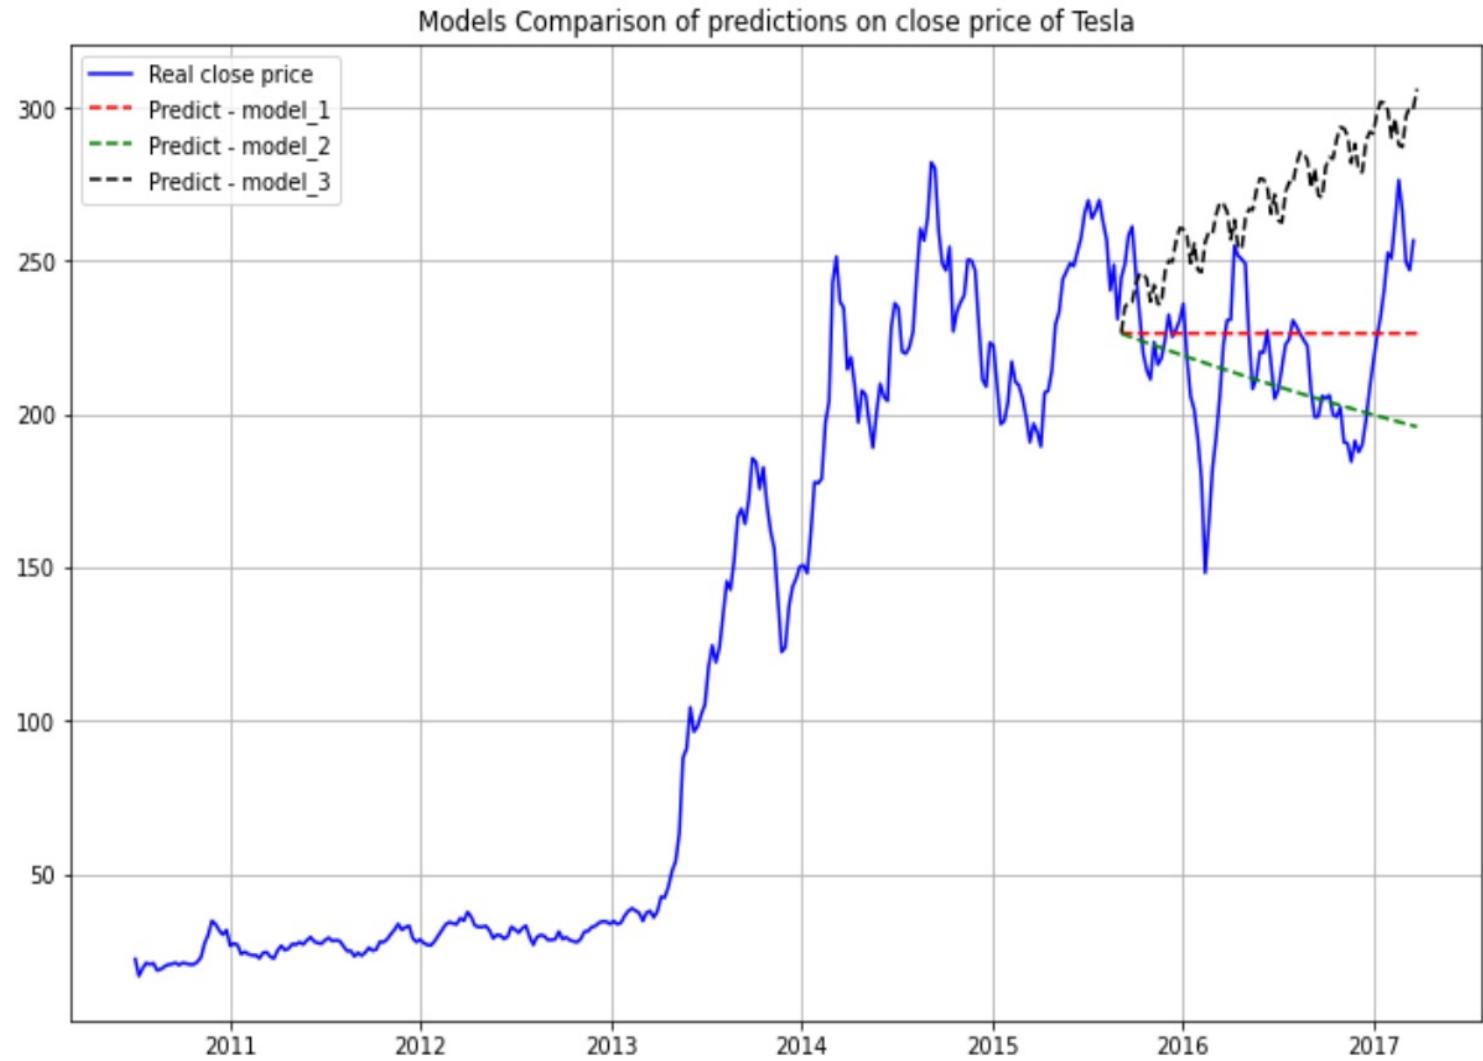
MSE of model_1: 50246.62748995037
MSE of model_2_best: 61624.472763307844
MSE of model_3_best: 280403.85610155656

Much worse...with higher mse.

Summary - key findings

- From model_3 prediction plot, it's clearly worse than other models, although some part of trends is similar.
- It shows that adding the seasonal effects in this dataset is misleading the predicted (increasing) trends which is far from real data.
- Comparing with model_1 and model_2, model_3 is much complicated but has much larger MSE.
- Model_1's order is much easier got by just checking acf / pacf than getting from pmd arima modeling (model_2, model_3).

In this case, from the comparing of ARIMA, pmd-ARIMA, pmd-SARIMA models, the close price prediction will prefer the model_1 ARIMA model as at least final choice (may not the best choice).



Insights and Suggestions



- The first date of stock price is at IPO date (2010-6-29), and normally a new stock will have big fluctuations in price after quite few periods. Tesla's stock prices stay stable (<\$50) for over 2 years.
 - Then, very huge increase with fluctuations come after 2013. It is obviously that, during 2013, there are business actions boosting its stock price, comparing with time sequence factors.
1. What if remove the first "stable" 2 years in the whole data? The train-set using 2013-2016 instead of 2010-2015? Will it affect on results of model_3?
 2. If no changes in original dataset, whether deep learning models will more accurate on predictions? I guess LSTM should get a try.

Follow-up 1 : change train data for SARIMA model



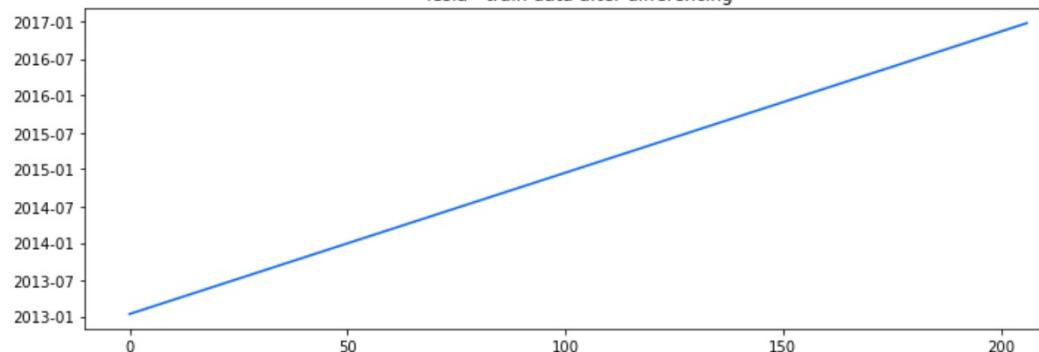
11.1.1 transfer to stationarity by differencing

```
1 new_train_diff = new_train.diff()
2 new_train_diff = new_train_diff.dropna()
```

```
1 plt.figure(figsize=(12,4))
2 plt.plot(new_train_diff.index)
3 plt.title("Tesla - train data after differencing")
```

```
Text(0.5, 1.0, 'Tesla - train data after differencing')
```

Tesla - train data after differencing



```
1 adfuller(new_train_diff)
```

```
(-11.120357537288317,
 3.4783714548659886e-20,
 0,
 206,
 {'1%': -3.4624988216864776,
 '5%': -2.8756749365852587,
 '10%': -2.5743041549627677},
 1428.0276199980033)
```

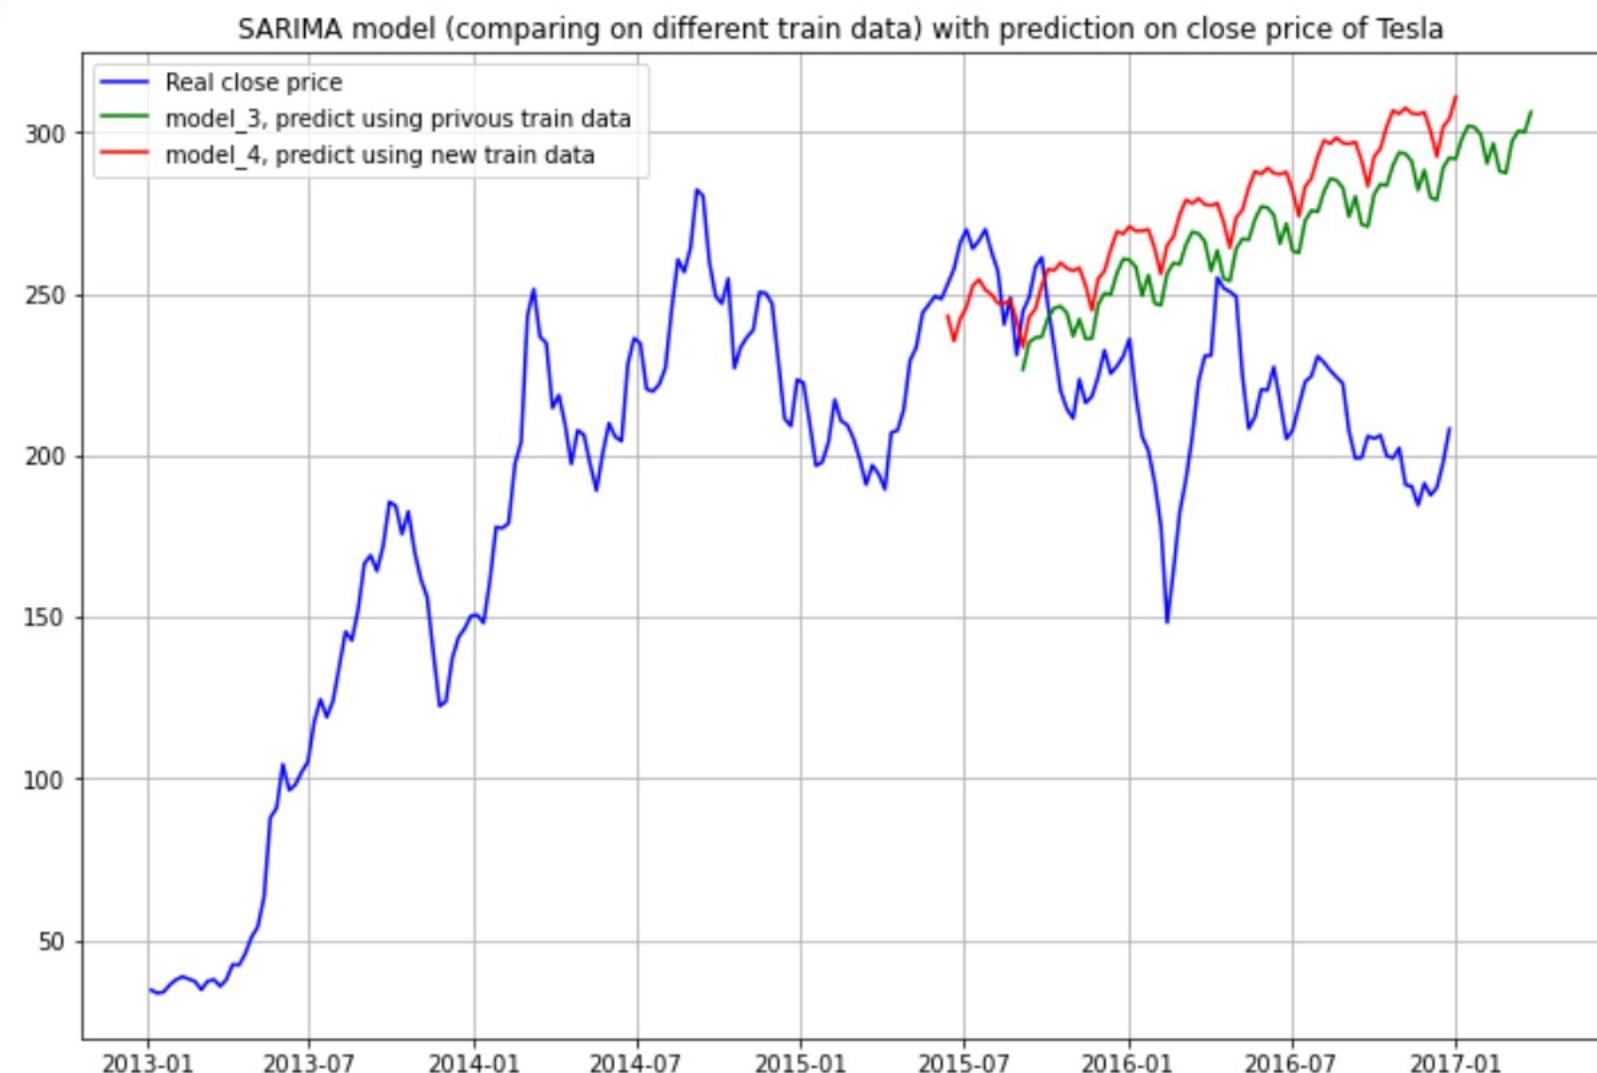
p-value < 0.05, after differencing, it is stationarity.

Model_4,
SARIMA,
order=(2,0,0),
Seasonal
order=(2,1,0,11)

previous MSE: 280403.85610155656
new Mse: 314807.81882714224

Changing train-data brings
No improving!

Previous sarima model order: (2, 0, 1)
Previous sarima model seasonal order: (2, 1, 0, 11)
New sarima model order: (2, 0, 0)
New sarima model seasonal order: (2, 1, 0, 11)



Follow-up 2: Deep learning LSTM model



Model preparing

- Data preprocessing (Min-max)
- Split train/ test dataset with x, y
- Transform to deep learning network data shape...

```
1 import math
2 train_data_len = math.ceil(len(data)*0.8)
3 train_data_len
```

1354

```
1 from sklearn.preprocessing import MinMaxScaler
2 scaler = MinMaxScaler(feature_range=(0,1))
3 trans_data = scaler.fit_transform(data)
4 trans_data

array([[0.02993635],
       [0.02971433],
       [0.02279455],
       ...,
       [0.88784039],
       [0.91122698],
       [0.9091918 ]])
```

```
1 train_data = trans_data[0: train_data_len,:]
2 train_x = []
3 train_y = []
4
5 for i in range(60,len(train_data)):
6     train_x.append(train_data[i - 60:i,0])
7     train_y.append(train_data[i,0])
8     if i<60:
9         print(train_x)
10        print(train_y)
11        print()
```

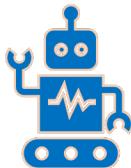
```
1 train_x, train_y = np.array(train_x), np.array(train_y)
2
3 train_x = np.reshape(train_x, (train_x.shape[0],train_x.shape[1],1))
4 train_x.shape
```

(1294, 60, 1)

```
1 test_data = trans_data[train_data_len-60:,:]
2 test_x = []
3 test_y = data[train_data_len:,:]
4 for i in range(60,len(test_data)):
5     test_x.append(test_data[i-60:i,0])
6 test_x = np.array(test_x)
7 test_x = np.reshape(test_x, (test_x.shape[0],test_x.shape[1],1))
8 test_x.shape
```

(338, 60, 1)

LSTM model building



```
1 model_5.summary()
```

Model: "sequential_3"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| <hr/> | | |
| lstm_4 (LSTM) | (None, 64) | 16896 |
| dense (Dense) | (None, 32) | 2080 |
| dense_1 (Dense) | (None, 1) | 33 |
| <hr/> | | |

Total params: 19,009

Trainable params: 19,009

Non-trainable params: 0

```
1 import tensorflow as tf
2 import keras
3 from keras.models import Sequential
4 from keras.layers import Dense, SimpleRNN, LSTM, Activation, Dropout
```

```
1 model_5 = Sequential()
2 model_5.add(LSTM(64,input_shape=(train_x.shape[1],1)))
3 model_5.add(Dense(32))
4 model_5.add(Dense(1))
5 model_5.compile(optimizer='adam', loss='mean_squared_error')
6 model_5.fit(train_x, train_y,batch_size=1, epochs=10)
```

Epoch 1/10

1294/1294 [=====] - 21s 15ms/step - loss: 0.0018

Epoch 2/10

1294/1294 [=====] - 19s 14ms/step - loss: 9.2596e-04

Epoch 3/10

1294/1294 [=====] - 18s 14ms/step - loss: 7.2786e-04

Epoch 4/10

1294/1294 [=====] - 19s 14ms/step - loss: 5.3970e-04

Epoch 5/10

1294/1294 [=====] - 19s 15ms/step - loss: 6.0790e-04

Epoch 6/10

1294/1294 [=====] - 19s 15ms/step - loss: 5.2898e-04

Epoch 7/10

1294/1294 [=====] - 19s 14ms/step - loss: 4.8845e-04

Epoch 8/10

1294/1294 [=====] - 17s 13ms/step - loss: 4.6672e-04

Epoch 9/10

1294/1294 [=====] - 17s 13ms/step - loss: 4.9823e-04

Epoch 10/10

1294/1294 [=====] - 16s 12ms/step - loss: 4.5805e-04

<keras.callbacks.History at 0x7fce154cecd0>

```
1 plt.figure(figsize=(20,12))
2 plt.title("LSTM model - Close stock price of Tesla")
3 plt.plot(train_df['Close'], 'b-')
4 plt.plot(valid_df[['Close','Predict']], linewidth = 4)
5 plt.xlabel("Time")
6 plt.ylabel("Close price")
7 plt.legend(['Train','Test','Predict'], loc='best');
```

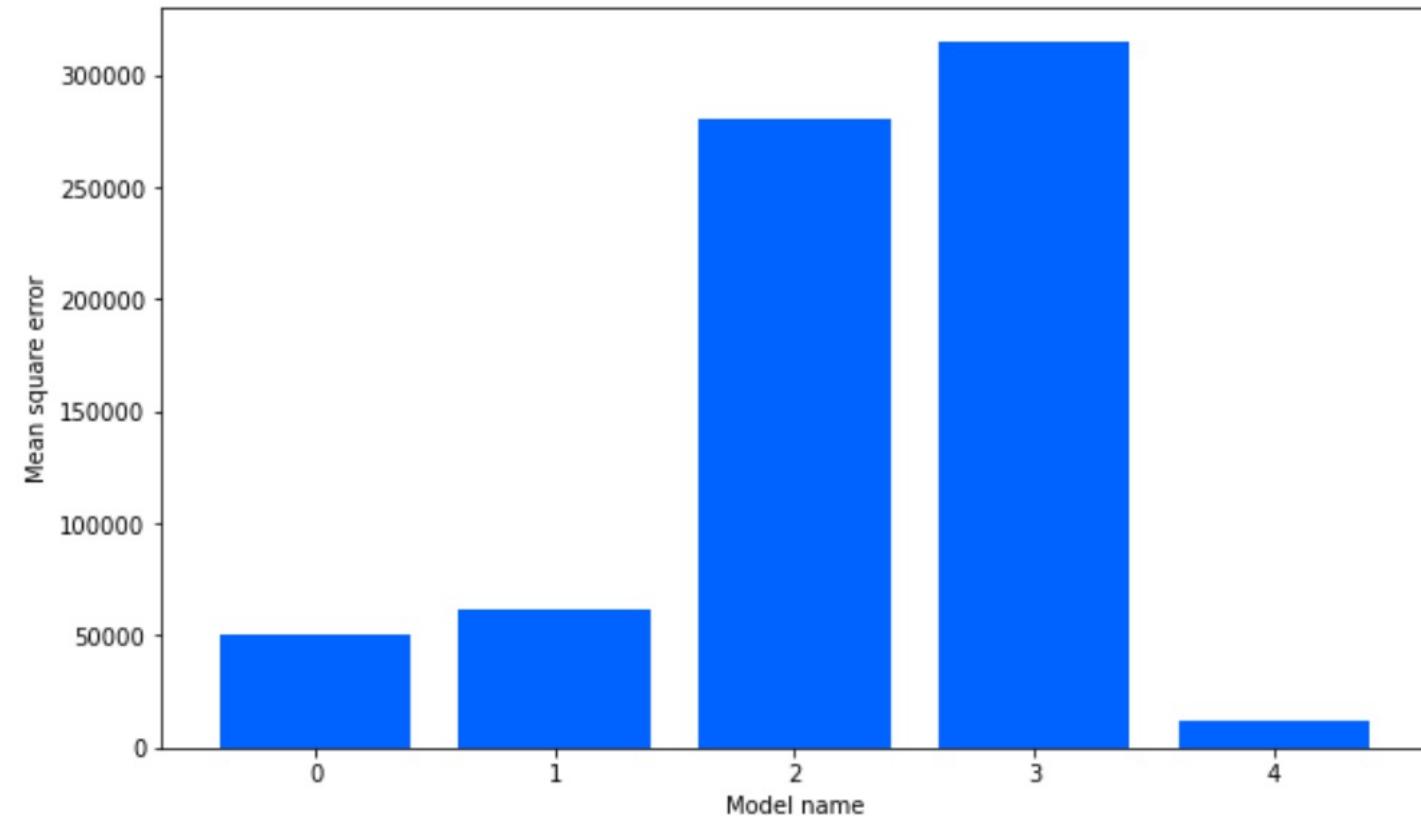
LSTM model prediction

Significantly
improved!



► Final conclusion

| | Model name | Mean square error |
|---|-------------------|-------------------|
| 0 | ARIMA model | 50,246.6 |
| 1 | PMD-ARIMA model | 61,624.5 |
| 2 | SARIMA model | 280,403.9 |
| 3 | (NEW)SARIMA model | 314,807.8 |
| 4 | LSTM model | 11,475.6 |



Conclusion:

Clearly, **LSTM** will be the best choice in prediction of stock price via 5 models comparing. Predictions(LSTM) are very similar to real data with the lowest MSE. Of course, as a disadvantage of deep learning models, it often occurs overfitting.

